

ExcelML Example01: Learning XOR End-to-End in Excel (Step-by-Step)

This walkthrough is written for **Excel-native users**. You will be told **exactly what values and formulas to enter into which cells**, and why each step matters.

You do *not* need prior machine learning knowledge.

By the end, you will:

- Train a real neural network inside Excel
 - See learning happen numerically and visually
 - Understand *why* XOR requires a hidden layer
 - Inspect the internal state of the model
-

Conceptual anchor (keep this in mind)

In ML terms, you are:

Training a neural network using gradient descent

In Excel terms, you are:

Using cells and formulas to explicitly control a learning process

Every formula you enter corresponds to a real ML operation.

Step 0 — Create a new worksheet

1. Open a **new blank workbook** in Excel
 2. Use the **first worksheet** (Sheet1)
 3. Ensure the ExcelML add-in is loaded
-

Step 1 — Enter the XOR dataset

We start by entering the smallest dataset that *cannot* be solved by a straight line.

1.1 Enter headers

Cell	Value
A1	X1
B1	X2
C1	y

1.2 Enter data

Cell	Value
A2	0
B2	0
C2	0
A3	0
B3	1
C3	1
A4	1
B4	0
C4	1
A5	1
B5	1
C5	0

At this point:

- A2:B5 is your input matrix X
- C2:C5 is your output vector y

Step 2 — Create a model

We now tell ExcelML what kind of neural network we want.

2.1 Choose a cell for the model ID

We will store the model ID in E2.

2.2 Enter the formula

In cell E2, type:

```
=DL.MODEL_CREATE("xor:in=2,hidden=8,out=1")
```

2.3 What you should see

- The cell will display a long string (a model ID)
- This string uniquely identifies a model stored in memory

2.4 Conceptual meaning

- `in=2` → two input features (X_1, X_2)
- `hidden=8` → one hidden layer with 8 neurons
- `out=1` → one output value (probability of class 1)

Excel is now holding a *reference* to a real neural network.

Step 3 — Create a training trigger

Excel recalculates formulas automatically. Training a neural network should **not** happen automatically.

We solve this with a trigger cell.

3.1 Choose a trigger cell

We will use Z1.

3.2 Enter an initial value

In cell Z1, type:

1

This value has no meaning except:

"Train when this value changes."

Step 4 — Train the model

This is the most important step.

4.1 Choose a cell for training output

We will use **E4**.

4.2 Enter the training formula

In cell **E4**, type:

```
=DL.TRAIN(E2, A2:B5, C2:C5, "epochs=1000,lr=0.1", $Z$1)
```

4.3 What this formula means

Argument	Meaning
E2	Model ID
A2:B5	Input data X
C2:C5	Target values y
epochs=1000	Number of training passes
lr=0.1	Learning rate
Z1	Training trigger

4.4 What you should see

- Training runs once
- The cell displays a small table with status, epochs, and final loss

If you **recalculate Excel**, training does *not* run again.

Step 5 — View the loss history

Loss tells us how wrong the model is.

5.1 Choose a cell

We will start the loss table in **E8**.

5.2 Enter the formula

In cell **E8**, type:

```
=DL.LOSS_HISTORY(E2)
```

5.3 What you should see

A spilled table like:

epoch	loss
1	~0.69
...	...
1000	much smaller

5.4 Optional: create a chart

1. Select the loss table
2. Insert → Line Chart

You are now watching a neural network learn.

Step 6 — Retrain on demand

6.1 Change the trigger

Change Z1 from 1 to 2.

6.2 What happens

- Training runs again
- Loss history resets
- New weights are learned

This mirrors real ML experimentation.

Step 7 — Inspect the model (advanced but illuminating)

You can look *inside* the model.

7.1 View weights

In an empty area, type:

```
=DL.WEIGHTS(E2, "fc1")
```

7.2 View gradients

```
=DL.GRADS(E2, "fc1")
```

7.3 View activations

```
=DL.ACTIVATIONS(E2, A2:B5, "fc1")
```

Each of these spills a table you can analyze or chart.

What you just did

Without writing code, you:

- Built a neural network
- Trained it using gradient descent

- Controlled training explicitly
 - Visualized learning
 - Inspected internal representations
-

Why XOR matters

XOR proves:

- Linear models are limited
- Hidden layers create expressive power

ExcelML lets you *see* this — not just read about it.

Next experiments

Try:

- Changing `hidden=2` , `hidden=16`
- Lowering or raising `lr`
- Removing the hidden layer

Each change teaches a real ML lesson.

Appendix: Conceptual and Pedagogical Deep Dive

This appendix expands on the walkthrough by explaining **what is really happening conceptually** at each stage, and *why the design choices in ExcelML matter pedagogically*.

It is intended for readers who want to move beyond “how” and into “why”.

A1. Why XOR is the perfect first neural network problem

XOR is not chosen because it is small — it is chosen because it is *revealing*.

Linear intuition breaks

A linear model tries to separate classes with a straight line (or plane). XOR cannot be separated that way:

- (0,0) and (1,1) belong to one class
- (0,1) and (1,0) belong to the other

No straight line can separate these two sets.

What XOR teaches

XOR demonstrates, in the smallest possible form:

- the limitation of linear models
- the necessity of hidden layers
- the idea of representation learning

When learners see XOR solved *only after adding a hidden layer*, the motivation for neural networks becomes concrete.

A2. What a neural network layer really is

A neural network layer performs three steps:

1. **Weighted sum of inputs**
2. **Add a bias**
3. **Apply a non-linear function**

In Excel terms:

- weights behave like coefficients
- bias behaves like an intercept
- the activation function bends space

Hidden layers are not magic — they are *coordinate transformations*.

A3. Why non-linearity matters

If you stack linear layers without non-linear activations, the result is still linear.

This is why:

- removing `Tanh` (or ReLU) causes XOR training to fail
- adding more linear layers does nothing

Non-linearity allows the model to:

- bend decision boundaries
- carve complex regions
- represent conditional logic

ExcelML makes this visible when you inspect activations.

A4. Gradient descent, intuitively

Gradient descent is often explained with calculus, but conceptually it is simple:

Take a small step in the direction that reduces error the most.

In ExcelML:

- each epoch is one full pass over the data
- the loss measures “how wrong” the model is
- gradients tell each weight which direction to move

By inspecting `DL.GRADS`, you can literally see:

- which weights are changing
 - which ones matter
-

A5. Why loss curves matter more than final accuracy

Beginners often focus on final results. Experts watch **learning curves**.

Loss curves tell you:

- whether learning is happening
- whether learning is stable

- whether the learning rate is too high or too low

In Excel, charts make these patterns visually obvious.

This is why ExcelML emphasizes `DL.LOSS_HISTORY` early.

A6. The role of the learning rate

The learning rate controls step size:

- Too small → learning is slow
- Too large → learning becomes unstable

Encouraged experiment:

- Try `lr=0.01`, `lr=0.1`, `lr=1.0`
- Observe how the loss curve changes

ExcelML turns hyperparameter tuning into a visual exercise.

A7. Why explicit training triggers matter pedagogically

Most ML tools hide training behind execution flow.

ExcelML makes training **explicit and intentional**:

```
=DL.TRAIN(..., trigger)
```

This teaches:

- training is a *decision*, not a side-effect
- models have state
- recomputation ≠ retraining

This mental model transfers directly to production ML systems.

A8. Inspectability vs black boxes

Traditional tools optimize for convenience.

ExcelML optimizes for:

- visibility
- traceability
- causal understanding

Being able to inspect:

- weights
- gradients
- activations

changes learning from *belief* to *observation*.

A9. Why Excel is uniquely suited for this

Excel offers:

- deterministic recalculation
- immediate feedback
- structured grids
- built-in visualization

These align unusually well with ML pedagogy.

ExcelML does not fight Excel — it leans into it.

A10. Transfer of learning

Nothing in ExcelML is toy-only.

The concepts learned here map directly to:

- PyTorch
- TensorFlow
- real-world ML pipelines

ExcelML is not a replacement — it is a **ramp**.

Closing thought

Most ML confusion comes from abstraction layers that hide causality.

ExcelML removes those layers.

Understanding comes from seeing cause and effect.

This appendix exists to help you build that intuition — deliberately, visually, and one cell at a time.